

# **mix-up : BEYOND EMPIRICAL RISK MINIMIZATION**

-Summary-

# Introduction

## Notation & Effect

- Example :  $(x, y) \sim P(X, Y)$ , where  $x$  = input,  $y$  = target
- Prediction :  $f(x)$ , where  $f \in \mathcal{F}$  is expected to satisfy  $f(X) = Y$
- Set of training data :  $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$ , where  $(x_i, y_i) \sim P$
- Expected risk :  $R(f) = \int l(f(x), y) dP(x, y)$   
(note : distribution  $P$  is unknown in practical situations)
- Empirical distribution :  $P_\delta(x, y) = \frac{1}{n} \sum_{i=1}^n \delta(x = x_i, y = y_i)$
- Empirical risk :  $R_\delta(f) = \int l(f(x), y) dP_\delta(x, y) = \frac{1}{n} \sum_{i=1}^n l(f(x_i), y_i)$
- Set of vicinity training data =  $\mathcal{D}_v = \{(\tilde{x}_i, \tilde{y}_i)\}_{i=1}^m$  (sampled  $m$  times from vicinity distribution)
- Gaussian vicinity distribution :  $v(\tilde{x}, \tilde{y} | x_i, y_i) = \mathcal{N}(\tilde{x} - x_i, \sigma^2) \delta(\tilde{y} = y_i)$  (Chapelle, 2000)
- Mix-up vicinity distribution :  $\mu(\tilde{x}, \tilde{y} | x_i, y_i) = \frac{1}{n} \sum_{j=1}^n \mathbb{E}_\lambda [\delta(\tilde{x} = \lambda x_i + (1 - \lambda) x_j, \tilde{y} = \lambda y_i + (1 - \lambda) y_j)]$ , where  $\lambda \sim \text{Beta}(\alpha, \alpha)$ ,  $\alpha \in (0, \infty)$
- Empirical vicinal risk :  $R_v(f) = \frac{1}{m} \sum_{i=1}^m l(f(\tilde{x}_i), \tilde{y}_i)$  (usually,  $m = n^2$ )

**Effect : Minimizing the empirical vicinal risk with mix-up improve test error (compared to ERM)**

**=> generalize better and robust to adversarial attack**

# Experimental results – Image classification

| Model             | Method                      | Epochs | Top-1 Error | Top-5 Error |
|-------------------|-----------------------------|--------|-------------|-------------|
| ResNet-50         | ERM (Goyal et al., 2017)    | 90     | 23.5        | -           |
|                   | <i>mixup</i> $\alpha = 0.2$ | 90     | <b>23.3</b> | <b>6.6</b>  |
| ResNet-101        | ERM (Goyal et al., 2017)    | 90     | 22.1        | -           |
|                   | <i>mixup</i> $\alpha = 0.2$ | 90     | <b>21.5</b> | <b>5.6</b>  |
| ResNeXt-101 32*4d | ERM (Xie et al., 2016)      | 100    | 21.2        | -           |
|                   | ERM                         | 90     | 21.2        | 5.6         |
|                   | <i>mixup</i> $\alpha = 0.4$ | 90     | <b>20.7</b> | <b>5.3</b>  |
| ResNeXt-101 64*4d | ERM (Xie et al., 2016)      | 100    | 20.4        | 5.3         |
|                   | <i>mixup</i> $\alpha = 0.4$ | 90     | <b>19.8</b> | <b>4.9</b>  |
| ResNet-50         | ERM                         | 200    | 23.6        | 7.0         |
|                   | <i>mixup</i> $\alpha = 0.2$ | 200    | <b>22.1</b> | <b>6.1</b>  |
| ResNet-101        | ERM                         | 200    | 22.0        | 6.1         |
|                   | <i>mixup</i> $\alpha = 0.2$ | 200    | <b>20.8</b> | <b>5.4</b>  |
| ResNeXt-101 32*4d | ERM                         | 200    | 21.3        | 5.9         |
|                   | <i>mixup</i> $\alpha = 0.4$ | 200    | <b>20.1</b> | <b>5.0</b>  |

## Note

- Model with higher capacities or longer training runs are the ones to benefit the most from mix-up method.

(ex : ResNet-50 (0.2%) <-> ResNeXt-101(0.5~ 0.6%), ResNet-50 at epoch 90 / 200)

# Experimental results – Speech data

| Model  | Method                          | Validation set | Test set    |
|--------|---------------------------------|----------------|-------------|
| LeNet  | ERM                             | <b>9.8</b>     | <b>10.3</b> |
|        | <i>mixup</i> ( $\alpha = 0.1$ ) | 10.1           | 10.8        |
|        | <i>mixup</i> ( $\alpha = 0.2$ ) | 10.2           | 11.3        |
| VGG-11 | ERM                             | 5.0            | 4.6         |
|        | <i>mixup</i> ( $\alpha = 0.1$ ) | 4.0            | 3.8         |
|        | <i>mixup</i> ( $\alpha = 0.2$ ) | <b>3.9</b>     | <b>3.4</b>  |

classification error

## Note

- Data preprocess :  
Data (65,000 utterances, ~1 sec, 30 classes) => Normalized spectrogram from waveforms  
=> Mix-up both at waveform and spectrogram levels => Train model (CNN-based)
- Similarly, model with higher capacity (=VGG-11) benefits the most from mix-up method.

# Experimental results – Memorization of corrupted labels

| Label corruption | Method   | Test error  |             | Training error |           |
|------------------|--|-------------|-------------|----------------|-----------|
|                  |  | Best        | Last        | Real           | Corrupted |
| 20%              | ERM  | 12.7        | 16.6        | 0.05           | 0.28      |
|                  | ERM + dropout ( $p = 0.7$ )                      | 8.8         | 10.4        | 5.26           | 83.55     |
|                  | <i>mixup</i> ( $\alpha = 8$ )                    | <b>5.9</b>  | 6.4         | 2.27           | 86.32     |
|                  | <i>mixup</i> + dropout ( $\alpha = 4, p = 0.1$ ) | 6.2         | <b>6.2</b>  | 1.92           | 85.02     |
| 50%              | ERM  | 18.8        | 44.6        | 0.26           | 0.64      |
|                  | ERM + dropout ( $p = 0.8$ )                      | 14.1        | 15.5        | 12.71          | 86.98     |
|                  | <i>mixup</i> ( $\alpha = 32$ )                   | 11.3        | 12.7        | 5.84           | 85.71     |
|                  | <i>mixup</i> + dropout ( $\alpha = 8, p = 0.3$ ) | <b>10.9</b> | <b>10.9</b> | 7.56           | 87.90     |
| 80%              | ERM  | 36.5        | 73.9        | 0.62           | 0.83      |
|                  | ERM + dropout ( $p = 0.8$ )                      | 30.9        | 35.1        | 29.84          | 86.37     |
|                  | <i>mixup</i> ( $\alpha = 32$ )                   | 25.3        | 30.9        | 18.92          | 85.44     |
|                  | <i>mixup</i> + dropout ( $\alpha = 8, p = 0.3$ ) | <b>24.0</b> | <b>24.8</b> | 19.70          | 87.67     |

## Note

- Certain portion of labels are replaced (corrupted) by random noises.
- Dropout + mix-up shows better performance on learning corrupted labels, compared to Dropout + ERM.

# Experimental results – Robustness to adversarial examples

test error

| Metric | Method       | FGSM        | I-FGSM |
|--------|--------------|-------------|--------|
| Top-1  | ERM          | 90.7        | 99.9   |
|        | <i>mixup</i> | <b>75.2</b> | 99.6   |
| Top-5  | ERM          | 63.1        | 93.4   |
|        | <i>mixup</i> | <b>49.1</b> | 95.8   |

(a) White box attacks.

| Metric | Method       | FGSM        | I-FGSM      |
|--------|--------------|-------------|-------------|
| Top-1  | ERM          | 57.0        | 57.3        |
|        | <i>mixup</i> | <b>46.0</b> | <b>40.9</b> |
| Top-5  | ERM          | 24.8        | 18.1        |
|        | <i>mixup</i> | <b>17.4</b> | <b>11.8</b> |

(b) Black box attacks.

## Note

- One of undesirable behavior of ERM is fragility to adversarial examples
- FGSM (Fast Gradient Sign Method) :  $x_{adv} = x + \epsilon * \text{sign}(\nabla_x l(f(x), y))$  (exploit SGD process)
- I-FGSM(Iterative-FGSM) :  $x_{adv}^0 = x, x_{adv}^{N+1} = X_N^{adv} + \alpha * \text{sign}(\nabla_x l(f(x), y))$
- Black box FGSM / I-FGSM implementation :  
Produce Adversarial examples (1<sup>st</sup> ERM model) ---> test robustness of 2<sup>nd</sup> / 3<sup>rd</sup> model  
(Use 3 models : 1<sup>st</sup> / 2<sup>nd</sup> = trained using ERM / 3<sup>rd</sup> = trained using mix-up)
- Mix-up method is better than ERM in terms of robustness to adversarial examples

# Experimental results – Non-image data

test error


| Dataset    | ERM  | <i>mixup</i> |
|------------|------|--------------|
| Abalone    | 74.0 | 73.6         |
| Arcene     | 57.6 | <b>48.0</b>  |
| Arrhythmia | 56.6 | <b>46.3</b>  |

| Dataset  | ERM  | <i>mixup</i> |
|----------|------|--------------|
| Htru2    | 2.0  | 2.0          |
| Iris     | 21.3 | <b>17.3</b>  |
| Phishing | 16.3 | 15.2         |

## Note

- Mix-up method also works well on non-image data (better than ERM, in general)

# Experimental results – Stabilization of GANs



iteration : 10, 100, 1000, 10000, 20000 / blue = dataset, orange = generated samples

## Note

- Original optimization of GAN :  $\max_g \min_d \mathbb{E}_{x,z} [l(d(x), 1) + l(d(g(z)), 0)]$   
( $d(\cdot)$  = discriminator,  $g(\cdot)$  = generator,  $l$  = binary cross entropy loss)
- Mix-up optimization of GAN :  $\max_g \min_{d,\lambda} \mathbb{E}_{x,z,\lambda} [l(d(\lambda x + (1 - \lambda)g(z)), \lambda)]$
- Training of mix-up GANs seems promisingly robust to hyper-parameter and architectural choices.



# Ablation study

AC : mix between all classes  
RP : mix between random pairs  
KNN : mix between k-nearest neighbors  
(Here,  $k = 200$ )

median test errors of the last 10 epochs

| Method   | Specification | Modified |        | Weight decay |                    |
|--|---------------|----------|--------|--------------|--------------------|
|  |               | Input    | Target | $10^{-4}$    | $5 \times 10^{-4}$ |
| ERM  |               | ✗        | ✗      | 5.53         | 5.18               |
| <i>mixup</i>                                       | AC + RP       | ✓        | ✓      | 4.24         | 4.68               |
|  | AC + KNN      | ✓        | ✓      | 4.98         | 5.26               |
| mix labels and latent representations<br>(AC + RP) | Layer 1       | ✓        | ✓      | 4.44         | 4.51               |
|  | Layer 2       | ✓        | ✓      | 4.56         | 4.61               |
|  | Layer 3       | ✓        | ✓      | 5.39         | 5.55               |
|  | Layer 4       | ✓        | ✓      | 5.95         | 5.43               |
|  | Layer 5       | ✓        | ✓      | 5.39         | 5.15               |

## Question : What is mix-up doing? (Ablation studies)

- Form of data augmentation that encourages the model  $f$  to behave linearly in-between training examples.
- ⇒ It turns out (experimentally) that mix labels + mix latent representations does not show better performance than mix labels + mix raw inputs (original mix-up)
- ⇒ As in experiments on speech data, it seems crucial to choose how to preprocess raw data and interpolate...  
(can't figure out best preprocess method in practice circumstances)

## Q: How to mix-up data in a good manner without background of data?

- ⇒ Mix labels + mix (trained) VAE outputs ?? (similar to mix labels + mix latent representations)